

### EDAF75 Database Technology

#### Lecture 3

Christian.Soderberg@cs.lth.se

January 28, 2020

#### Today

- ▶ We'll first continue with some SQL
- ▶ Then we'll talk about how to make a model of a system, and then turn that model into a database

- ▶ We need course representatives
- ▶ You can now sign up for the labs (see web site)
- ▶ If you want a QA-session on Friday, please sign up on Moodle
- ▶ This week we'll continue with SQL queries, and look at *database modeling*
- ▶ Lab 1 is next week, it's an exercise in SQL queries
- ▶ Lab 2 is the week after next, and deals with modeling (which we'll talk about this week)

#### Short recapitulation

- ▶ A *relational database* is a collection of one or more *tables*, where each table has a fixed set of *columns* and a varying number of rows – *all cells contain primitive values*
- ▶ Simple queries (SELECT-FROM-WHERE)
- ▶ Set operations (UNION, INTERSECT, EXCEPT)
- ▶ Simple functions and aggregate functions
- ▶ Grouping (GROUP BY-HAVING)
- ▶ Subqueries
- ▶ Joins (CROSS-, INNER-, LEFT OUTER-, ...)


## Today

- ▶ Views: 'virtual' tables which enables us to simplify queries (in SQLite it's essentially a stored SELECT-statement)
- ▶ Common Table Expressions (CTE): a view which exists only for the duration of a single query
- ▶ Database modeling



## Elements of an E/R Model

- ▶ **Entity Sets:** these are the 'objects' of our model, they correspond to classes in a traditional object oriented model
- ▶ **Attributes:** properties of our objects – must be primitive values (see the next slide)
- ▶ **Relationships:** associations between our entity sets
- ▶ We will typically convert entity sets to tables (*relations*), and attributes to columns in our tables – relationships will be dealt with according to their multiplicities



## Modeling

- ▶ To design a database, we'll start out with what's called an Entity/Relation Model (E/R Model)
- ▶ There are many 'standards' for drawing E/R diagrams, we'll use UML class diagrams – it's becoming increasingly popular for database modeling



## Primitive values in our models

- ▶ integers: INT, INTEGER, ...
- ▶ real numbers: REAL, DECIMAL(w, d), ...
- ▶ strings: TEXT, CHAR(n), VARCHAR(n)
- ▶ dates: DATE, TIME, TIMESTAMP, ...
- ▶ blobs (binary large objects): BLOB (only in some databases)



## UML class diagrams

- ▶ We'll use UML classes in approximately the same way as you may have seen in earlier courses, with some caveats:
  - ▶ There will be no methods in our classes
  - ▶ All our attributes will be primitive and public
  - ▶ We won't bother much with aggregates and compositions, we'll use plain associations instead
  - ▶ We'll be very careful to mark multiplicities everywhere
  - ▶ We will be preoccupied by what constitutes a key for each entity set

## Example

**Exercise:** Create a simple model for students taking courses at LTH

## Example

**Exercise:** Create a model for students taking courses at LTH, be a little bit more elaborate this time

## Keys

- ▶ A *superkey* is a set of attributes for which there are no tuples (rows) that have the same values for the attributes of the set
- ▶ A *key* is a *minimal* set of attributes which uniquely identifies each row in an entity set
- ▶ Sometimes there are several candidate keys, when we model our database we pick one of them and call it our *primary key*
- ▶ We need keys to find values easily, and to keep track of related rows in other tables

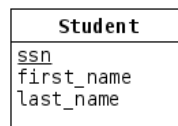
## Natural keys and invented keys

- ▶ Sometimes keys occur naturally in the problem domain, we call them *natural keys* or *business keys*
- ▶ Some entity sets have no natural keys, we call them *weak entity sets*
- ▶ Sometimes we invent keys by introducing some artificial attribute, these keys are called *invented keys*, *surrogate keys*, *synthetic keys*, ...

## Natural keys and invented keys


- ▶ Adding an invented key to a "weak" entity set makes it a regular entity set
- ▶ Whether to use an invented key or not is often a question of simplicity vs efficiency:
  - ▶ Without an invented key we sometimes get an unwieldy key (several attributes)
  - ▶ With an invented key our tables and queries only need a single key column, *but finding the key may require additional joins*
- ▶ If an attribute might change over time, it's not a good choice for a key – it would require us to update all tables which uses the old value

## UML class diagrams – classes



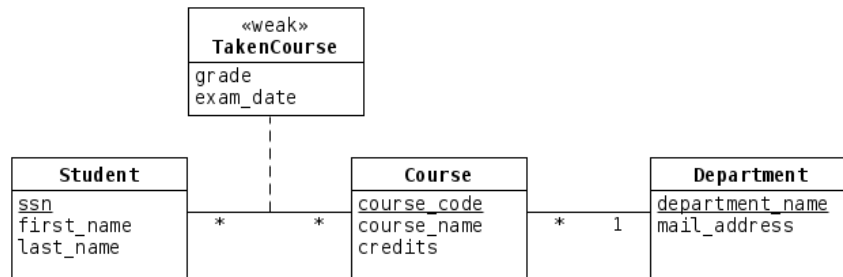
- ▶ The entity/class name in singular
- ▶ Only one box (since we have no methods)
- ▶ We will underline keys

## UML class diagrams – associations



- ▶ We always mark multiplicity on our associations
- ▶ We use associations instead of attributes whenever the value is a reference to another object

## UML class diagrams – association classes



- ▶ Sometimes the association between two entity sets contains data itself
- ▶ We use an *association class* to capture that data
- ▶ Often we can use either an association class between two entity sets, or another entity set 'between' them

## Foreign keys

- ▶ A foreign key is an attribute which refers to a key in another relation
- ▶ If we declare a column to be a FOREIGN KEY, our database will check that there is a corresponding value in the referenced table

## Translating E/R models to databases

- ▶ For each entity set (class) we create a relation (table) with the same attributes (columns) as the entity set
- ▶ For relationships (associations), what we do depend on their multiplicity:
  - ▶ '1 – 1'-associations often means that the two entity sets can be merged into one entity set
  - ▶ '\* – 1'-associations will be turned into *foreign keys* in the tables on the \*-side
  - ▶ '\* – \*'-associations will be turned into relations (tables)
- ▶ Inheritance and other multiplicities require more consideration (we'll look at it next time)