

EDAF75  
Database Technology

Lecture 4

Christian.Soderberg@cs.lth.se

January 30, 2020

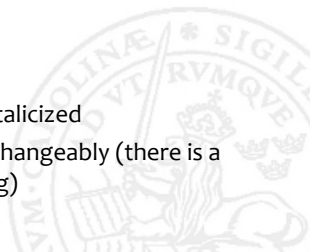


- ▶ Solve the 'average-highschool-size' problem with a subquery instead of a view
- ▶ Solve the 'exactly-two-chemistry-laureates-when-summer-olympics-in-Europe' problem without a subquery (but use the CTE's we wrote)



Database design

- ▶ To design a database, we first create a model which captures:
  - ▶ objects (entity sets),
  - ▶ the properties of the objects (attributes)
  - ▶ associations (relationships) between our objects – we're meticulous when it comes to multiplicities
- ▶ We then translate the model into a set of *relations* – these relations will become tables in our database
- ▶ We write relations as:  
`cars(license_plate, brand_id, model)`  
where primary keys are underlined, and foreign keys are italicized
- ▶ We'll sometimes use the terms 'relation' and 'table' interchangeably (there is a 1 – 1 correspondence, but they're not really the same thing)



Exercise

At LTH, courses are given every year, and each year a number of students register to take courses. Describe this in the simplest possible ER-model.




## Translation from model to relations

- ▶ Classes (entity sets) will become relations – we use *snake\_case* and plural form for the relations
- ▶ Attributes of our classes will become attributes of our relations
- ▶ Associations are handled according to multiplicity

## Translation from model to relations – example

We have two entity sets, Book and Author, and an association between them:



Create relations for them if:

- ▶ each book is written by one author
- ▶ a book can be written by several authors

## Translation from model to relations – example

Define relations based on the model of students taking courses at LTH which we just developed.

## Translation of associations

- ▶ \* – 1 associations are translated into *foreign keys* on the \* side
- ▶ \* – \* associations are translated into relations with foreign keys referencing both sides
- ▶ \* – 0..1 associations are a bit special:
  - ▶ if it's mostly 1 on the right side, we can use the first method above, and use NULL where we have no associated object
  - ▶ otherwise we can use the second method above (a new relation with two foreign keys)
- ▶ Other multiplicities require some handiwork

## Translating 0..1 – 0..1 associations – example

**Exercise:** We want to keep track of people and dogs, and assume a person can only own one dog, and that a dog can be owned by at most one person.

What relations do we use if:

- ▶ Almost all dogs have an owner
- ▶ Almost every person have a dog
- ▶ Only some people own dogs, and many dogs are without an owner

## Translating 0..1 – 0..1 associations

- ▶ If almost all dogs have owners:

```
people(ssn, ...)  
dogs(id, ..., owner_ssn)
```

- ▶ If almost everyone own a dog:

```
people(ssn, ..., dog_id)  
dogs(id, ...)
```


- ▶ If only some people own dogs, and many dogs are without an owner:

```
people(ssn, ...)  
dogs(id, ...)  
dog_ownerships(owner_ssn, dog_id)
```

## Translation of association classes


- ▶ For \* – \*-relations:
  - ▶ the \* – \*-relation itself will give us a new relation
  - ▶ the attributes of the association class will be attributes of this relation
- ▶ For \* – 1-relations:
  - ▶ the attributes of the association class will end up together with the foreign key on the \*-side

## Translating inheritance into relations



- ▶ Create one relation for each class (entity set), and reference from subclasses to superclasses using foreign keys
- ▶ Create relations only for concrete classes (entity sets)
- ▶ Create one big relation, with all possible attributes (with a lot of NULL values)


## Translating inheritance into relations



- Create one relation for each class (entity set), and reference from subclasses to superclasses using foreign keys:

```
people(ssn, name, ...)
students(ssn, stil_id, ...)
conductors(ssn, orchestra, ...)
```


## Translating inheritance into relations



- Create relations only for concrete classes (entity sets):

```
students(ssn, name, stil_id, ...)
conductors(ssn, name, orchestra, ...)
```

## Translating inheritance into relations



- Create one big relation, with all possible attributes (with a lot of NULL values)

```
people(ssn, name, stil_id, orchestra, ...)
```

## Defining tables in SQL

- To create a table, we use the CREATE TABLE command (and an optional DROP TABLE first):

```
DROP TABLE IF EXISTS books;
```

```
CREATE TABLE books (  
    isbn    TEXT,  
    title   TEXT,  
    year    INT,  
    ...  
);
```

- For each row we define a type (see next slide)
- For each row we can add constraints

## Types in our tables

- ▶ integers: INT, INTEGER, ...
- ▶ real numbers: REAL, DECIMAL(w,d), ...
- ▶ strings: TEXT, CHAR(n), VARCHAR(n)
- ▶ dates: DATE, TIME, TIMESTAMP, ...
- ▶ blobs (binary large objects): BLOB (only in some databases)

## Primary keys in table definitions

- ▶ We can declare an attribute to be primary key 'in place':

```
CREATE TABLE books (  
  isbn    TEXT PRIMARY KEY,  
  title   TEXT,  
  year    INT,  
  ...  
);
```

- ▶ We can also declare it separately (especially useful when the key contains several attributes):

```
CREATE TABLE books (  
  isbn    TEXT,  
  title   TEXT,  
  year    INT,  
  ...  
  PRIMARY KEY (isbn);  
);
```

## Foreign keys in table definitions

- ▶ We can declare foreign keys 'in place':

```
CREATE TABLE books (  
  isbn    TEXT PRIMARY KEY,  
  title   TEXT,  
  author  TEXT REFERENCES authors(author),  
  ...  
);
```

- ▶ We can also declare it separately:

```
CREATE TABLE books (  
  isbn    TEXT,  
  title   TEXT,  
  author  TEXT,  
  ...  
  FOREIGN KEY (author) REFERENCES authors(author);  
);
```

## Exercise

Create tables for the "students and courses at LTH" database which we designed earlier

## Exercise, from Tuesday

Solve the first part of problem 1 on the exam from april 2017



## Exercise

Create tables for the database we just designed



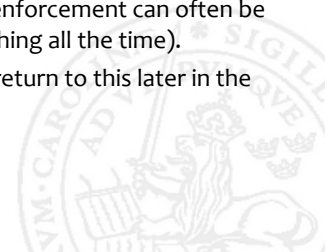
## Exercise

Create relations for the model we just developed



## Some other constraints

- ▶ We can declare a column to be:
  - ▶ NOT NULL
  - ▶ UNIQUE
  - ▶ DEFAULT <value>
  - ▶ CHECK <condition>
- ▶ These properties are enforced by the database, but the enforcement can often be temporarily turned off (it does take time to check everything all the time).
- ▶ We can also define *triggers* to enforce constraints, we'll return to this later in the course



## Inserting values

- ▶ We can insert values using INSERT:

```
INSERT
INTO  students(s_id, s_name, gpa, size_hs)
VALUES (123, 'Amy', 3.9, 1000),
       (234, 'Bob', 3.6, 1500),
       ...
```

- ▶ We don't have to provide values for columns with default values
- ▶ We also don't have to provide values for primary keys which are declared as INTEGER – they will get a new unique integral value (hence the moniker *database sequence number*)

## Updating values

- ▶ We can update values using UPDATE:

```
UPDATE students
SET    gpa = min(1.1 * gpa, 4.0)
WHERE  s_name LIKE 'A%';
```

- ▶ All rows are updated if we don't provide a WHERE clause

## Deleting values

- ▶ We can delete values using DELETE:

```
DELETE
FROM  applications
WHERE s_id = 123
```

- ▶ Beware that the innocent looking:

```
DELETE
FROM  applications
```

empties the whole table

## Variants

- ▶ There are various variants of the INSERT and UPDATE commands, such as:

- ▶ INSERT OR REPLACE
- ▶ INSERT OR IGNORE
- ▶ INSERT OR FAIL
- ▶ INSERT OR ROLLBACK
- ▶ UPDATE OR REPLACE
- ▶ UPDATE OR IGNORE
- ▶ UPDATE OR FAIL
- ▶ UPDATE OR ROLLBACK

- ▶ They are useful when an insertion or update would break some constraint

## Keys, superkeys, and invented keys

- ▶ A key is a set of attributes which uniquely identifies each row in a table
- ▶ In our college application database, `s_id` is unique in the `students` table, but `s_name` isn't
- ▶ Of course, the tuple (`s_id`, `s_name`) is also unique, but here `s_name` is redundant, so we call (`s_id`, `s_name`) a *superkey* (which, despite its name, is way less fancy than a regular key)
- ▶ The `s_id` is an artificial value which we associate with each student, and we call it an *invented key* (or *surrogate key*)

## Natural keys, and compound keys

- ▶ In the `colleges` table, `c_name` is unique in our small example, so we can use it as a key
- ▶ It is a value which occurs naturally in our problem domain, so we call it a *natural key*
- ▶ If we had a database with more colleges, we would soon end up with universities with the same name, but in different states – in that case we could use (`c_name`, `state`) as a key, it would be a *compound key*

## Foreign keys

- ▶ In the `applications` table, we have two attributes, `s_id` and `c_name` which refers to keys in other tables, they are called *foreign keys*
- ▶ The key for the `applications` table is (`s_id`, `c_name`, `major`)

## More about invented keys

- ▶ If a key ever changes, we may have to update in many places in our database – so we normally avoid natural keys which might change, and use an invented key instead
- ▶ Having a compound key can be a bit unwieldy, it is sometimes better to replace it with a simple invented key
- ▶ On the other hand: having an invented key requires us to do more joins



## Different kinds of invented keys

- ▶ We can use an increasing sequence of integers as invented key – it is very common, but it's predictable, and reveals something about the state of the database
- ▶ A *uuid* (universally unique identifier) is a  $\approx 128$ -bit random number, and it's a good choice for an invented key – we can safely assume it will be unique, and it doesn't reveal anything about the state of our system
- ▶ In SQLite3 we can use `lower(hex(randomblob(16)))` to create something akin to a *uuid* (but it takes up more memory)

## Keys, technically

- ▶ A key is a (minimal) set of attributes which uniquely identifies each row in a table
- ▶ For some relations we don't care about keys (e.g., a log of events)
- ▶ For some relations we may have several possible (*candidate*) keys – we then choose one of them as primary key
- ▶ For some relations there is no key using their own attributes, we call the corresponding entity set *weak*
- ▶ For a weak entity set, we can often use keys in associated entity sets (called *supporting entity sets*) to uniquely define a row, or just use an invented key

## Generating invented keys

- ▶ In SQLite3 we can get a *uuid*-lookalike using:

```
CREATE TABLE students (  
  s_id      TEXT DEFAULT (lower(hex(randomblob(16)))),  
  s_name    TEXT,  
  gpa       DECIMAL(3,1),  
  size_hs   INT,  
  PRIMARY KEY (s_id);  
);
```

- ▶ We'll soon see how the `DEFAULT` clause is used
- ▶ The database doesn't have to check if the generated value is unique, since the chance of a collision is ridiculously low
- ▶ The most recent version of SQLite3 (Sqlite 3.31) has a `uuid()`-function