

EDAF75 Database Technology

Lecture 5

Christian.Soderberg@cs.lth.se

February 4, 2020



- ▶ Lab 1 this week, lab 2 and lab 3 the following two weeks
- ▶ The Thursday lab sessions will be *packed*, so bring something else to do while you wait




Today's lecture

- ▶ A few words about keeping track of state
- ▶ Using SQLite3 from the command line
- ▶ Connecting to a database from Java using JDBC
- ▶ Connecting to a database from Python using the `sqlite3` package



Handling state

- ▶ How do we keep track of the balance of a bank account?
- ▶ Two possible solutions:
 - ▶ using a mutable attribute (balance)
 - ▶ by saving all updates (i.e., deposits and withdrawals)
- ▶ Not having a mutable balance requires us to do additional computations, but it brings some really nice benefits:
 - ▶ We can now explain why we have the current balance
 - ▶ We can see what the balance has been at different points of time in the past
 - ▶ We don't have to lock the database while updating the balance
- ▶ Saving changes instead of state is sometimes called *event sourcing*



Running SQLite3

- ▶ DBMS's such as PostgreSQL, MariaDB, Oracle, MySQL, etc., run as servers, and offer clients which talk to the servers
- ▶ SQLite3 doesn't run as a server, but it has a simple text based client which lets us manipulate its databases
- ▶ There are also some nice GUIs for SQLite3, just as there are for most DBMS's

SQLite3– command line client

Some useful SQLite3 commands:

- ▶ `.help`: makes this slide futile
- ▶ `.tables`: shows all tables in the current database
- ▶ `.schema <table>`: shows how a table is defined
- ▶ `.dump <table>`: gives INSERT statements for creating the specified table
- ▶ `.import <filename> <table>`: imports data into a table
- ▶ `.read <filename>`: reads a script from a file
- ▶ `.save <filename>`: writes the current database to file

SQLite3 format

We can get various output formats, using `.mode` – some examples (there are more):

- ▶ `csv`: Comma-separated values
- ▶ `column`: Left-aligned columns
- ▶ `html`: HTML code
- ▶ `insert`: SQL insert statements
- ▶ `line`: One value per line
- ▶ `list`: Values delimited by some separator
- ▶ `tabs`: Tab-separated values

Java Database Connectivity (JDBC)

- ▶ Standard classes for handling database connections
- ▶ Can handle all relevant relational databases
- ▶ Based upon *connections*, *statements*, and *result sets*
- ▶ Lots of things can go wrong when we connect to databases, so JDBC requires lots of exception handling
- ▶ There are also some alternative, non-standard libraries, such as `sql2o` (but they often depend on JDBC)

Connection

- ▶ Used to set up a session to a database (in the case of SQLite3, we don't really need to connect, the database is in a file on our hard drive, or even in memory)
- ▶ Creates Statement-objects, which we can use to send SQL statements to our database
- ▶ Connections also handle *transactions* (we'll talk about that later in the course)
- ▶ Used to require some strange incantations to set up, but no longer does

Statements

- ▶ Statement: a simple but unsafe kind of statement (amenable to *SQL injection*)
- ▶ PreparedStatement: a precompiled statement, safer, and more efficient when executed multiple times
- ▶ We *always* use 'try-with-resources' to create statements, to make sure that they are closed properly when we finish

PreparedStatement

- ▶ Created with `prepareStatement(str)` on a connection, where `str` is a `String` containing a query or statement
- ▶ The query or statement can contain parameters, marked as `?` – they get their values with various `setXXX`-methods
- ▶ We typically call `boolean execute()` if we have an update statement
- ▶ For queries, we call `ResultSet executeQuery()`

ResultSet

- ▶ A `ResultSet` which represents the table of data returned from an SQL query
- ▶ It's a kind of iterator (but doesn't implement any iterator interface), we call `next()` to jump to the next row, and it returns `false` if there is no next row
- ▶ We can use various `getXXX`-methods to fetch data, both positionally and by name

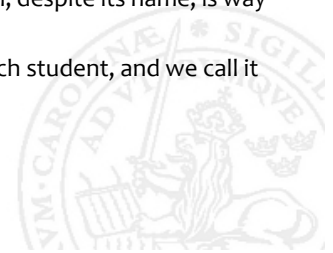
JDBC, code sample

```
var found = new ArrayList<ReturnType>( );
var query =
    "SELECT ... \n" +
    "FROM ... \n" +
    "WHERE ... = ? \n" +
    "... \n";
try (var ps = conn.prepareStatement(query)) {
    ps.setString(1, ...); // set parameter value
    var rs = ps.executeQuery( );
    while (rs.next()) {
        found.add(ReturnType.fromRS(rs));
    }
} catch (SQLException e) {
    e.printStackTrace( );
}
return found;
```



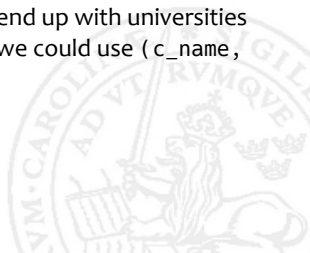
Keys, superkeys, and invented keys

- ▶ A key is a set of attributes which uniquely identifies each row in a table
- ▶ In our college application database, `s_id` is unique in the students table, but `s_name` isn't
- ▶ Of course, the tuple (`s_id`, `s_name`) is also unique, but here `s_name` is redundant, so we call (`s_id`, `s_name`) a *superkey* (which, despite its name, is way less fancy than a regular key)
- ▶ The `s_id` is an artificial value which we associate with each student, and we call it an *invented key* (or *surrogate key*)



Natural keys, and composite keys

- ▶ In the colleges table, `c_name` is unique in our small example, so we can use it as a key
- ▶ It is a value which occurs naturally in our problem domain, so we call it a *natural key*
- ▶ If we had a database with more colleges, we would soon end up with universities with the same name, but in different states – in that case we could use (`c_name`, `state`) as a key, it would be a *composite key*



Foreign keys

- ▶ In the applications table, we have two attributes, `s_id` and `c_name` which refers to keys in other tables, they are called *foreign keys*
- ▶ The key for the applications table is (`s_id`, `c_name`, `major`), it is a composite key
- ▶ If we needed to refer to the applications table from other tables, we'd need the composite key as a foreign key



More about invented keys

- ▶ To avoid having big composite keys being dragged around our databases, we often create invented keys
- ▶ If the value of a key ever changes, we may have to update in many places in our database – so we normally avoid natural keys which might change, and use an invented key instead
- ▶ On the other hand: having an invented key requires us to do more joins

Different kinds of invented keys

- ▶ We can use an increasing sequence of integers as invented key – it is very common, but it's predictable, and reveals something about the state of the database
- ▶ A *uuid* (universally unique identifier) is a ≈ 128 -bit random number, and it's a good choice for an invented key – we can safely assume it will be unique, and it doesn't reveal anything about the state of our system
- ▶ In SQLite3 we can use `lower(hex(randomblob(16)))` to create something akin to a uuid (but it takes up more memory)
- ▶ The most recent version of SQLite3 (Sqlite 3.31) has a `uuid()`-function

Generating invented keys

- ▶ In SQLite3 we can get a uuid-lookalike using:

```
CREATE TABLE students (  
  s_id      TEXT DEFAULT (lower(hex(randomblob(16)))),  
  s_name    TEXT,  
  gpa       DECIMAL(3,1),  
  size_hs   INT,  
  PRIMARY KEY (s_id);  
);
```

- ▶ The database doesn't have to check if the generated value is unique, since the chance of a collision is ridiculously low