

EDAF75
Database Technology

Lecture 10

Feb 27, 2020

Today

- ▶ Short intro to the project
- ▶ More on foreign keys
- ▶ SQL injection
- ▶ Indexes and query planning
- ▶ Limitations of SQL
- ▶ Stored procedures
- ▶ Alternatives to SQL (NoSQL)
- ▶ Scaling and Map/Reduce

Administration

- ▶ On Tuesday next week, March 3, we'll have a peer review session where you'll discuss your project designs with other groups – *participation is required from everyone who wants credits for the project*, and you must sign up before 23:59 on Sunday
- ▶ To report the progress of your project, you have to use a git repository on `gitlab.com`, please let me know if you want to learn how to use git
- ▶ Today we'll have a quick look at the project, and then go through the remainder of the course material

Project

- ▶ Design a database for a small bakery (Krusty)
- ▶ Participate in a design meeting, where your group will discuss your design with two other groups – you'll also give feedback on the other groups designs
- ▶ You must bring an ER-diagram showing your design to the meeting, in three copies (one for each group) – the design doesn't have to be perfect, but should be reasonably good (so, spend at least a couple of hours on it beforehand)
- ▶ The ER-diagram *must* be in proper UML (i.e., no "Crows foot", "Chen", or other notations)

Project

- ▶ Each group will 'report' to a project supervisor (PS)
- ▶ Your code and report should be kept in a *private* repository on `gitlab.com`, and you'll share it with your PS
- ▶ The report should be a Markdown file called `README.md` – the problem text describes what should be in it (we try to keep things simple, so it isn't much)

Project

- ▶ As soon as you're finished, push your code and report to `gitlab.com`, and notify your TA
- ▶ Each project supervisor will have *lots* of project to keep track of, and they will at the same time be busy with other courses, so unfortunately the feedback may sometimes take several days
- ▶ Deadline: you must have pushed your final (approvable) version to `gitlab.com` no later than 23:59 on April 30

Project

- ▶ Some of you have programmed a lot before taking this course, many of you haven't – we'll require that all of you do your best (but no more), *and that you follow common coding guidelines*, such as:
 - ▶ Never, ever, use tabs in your Java or Python code!
 - ▶ Always use correct indentation (4 spaces, in both Java and Python), and place braces where the standard guidelines put them
 - ▶ Functions/methods should generally do only one thing each – a function/method which does many different things should be broken into several functions/methods
 - ▶ Use proper names – functions/methods/parameters should have descriptive names, local variables can be shortened (i.e., how descriptive a name should be depends on how big its scope is)

Transactions and triggers

- ▶ Last time we looked into transactions and triggers
- ▶ The project will be much easier to implement if you use transactions and triggers properly!
- ▶ See some examples of transactions and triggers in the notebook

Example

Example

Write a program which keeps track of the names of your friends, and write code to insert new friends

SQL Injection

- ▶ We must be careful before putting user data into our database – never concatenate parameters into a query!
- ▶ SQL injection is when a user sends a string which alters the intended meaning of our SQL statement
- ▶ A classic example is the statement

```
stmt = "SELECT * FROM users WHERE name = '"  
      + user_name + "'";"
```

and user_name gets the value "" OR '1'='1"

- ▶ Using PreparedStatement instead of Statement makes our code safer
- ▶ In the project, we'll require that you use a PreparedStatement where a Statement would have been dangerous

SQL injection

- ▶ On wikipedia there is a long list of known SQL injections (and we probably don't hear about most of the successful ones)
- ▶ So, SQL injection is a real thing, and we'd better be safe than sorry – using a PreparedStatement is a very simple protective measure
- ▶ In languages such as C/C++, there are relatives to SQL injection – one well known example is "buffer overruns", in which someone crafts a message which overflows the buffer created to hold the reply (this was one of the exploits used by the "Morris worm")
- ▶ Very few programmers have regretted writing code which was 'too safe'

Indices

- ▶ When we define a primary key, the database creates a special *index* to make searches for the key fast – indexes can also speed up joins and ORDER BY statements
- ▶ We can create our own indexes, with as many columns as we want

```
CREATE INDEX names_and_ages  
ON employees(last_name, first_name, birth_year);
```

- ▶ This index will speed up searches for:
 - ▶ last name
 - ▶ last name and then first name
 - ▶ last name and birth year (with some fiddling)
- ▶ The index above will not help much if we're just searching for first name, or birth year
- ▶ Indices make some things faster, but insertions and deletions will become slower
- ▶ We can sometimes create a *covering index*, it is an index which includes the value we're normally looking for – using a covering index the DBMS will not even have to look at the table itself

Indices and the Query Planner

- ▶ For each SQL statement, there might be thousands of ways to perform the operation
- ▶ Before a DBMS executes a statement, its *query planner* takes a good look at it, to find a way to execute the statement efficiently
- ▶ Indices are very important during the planning – when searching for a value which isn't indexed properly, the DBMS might have to do a linear search through a table
- ▶ Sometimes big joins (especially cross joins) gives the query planner so many alternatives that the planning itself takes substantial time
- ▶ In SQLite3 there is a command `EXPLAIN QUERY PLAN` which explains what will happen during a given query

Exercise

Exercise

Define a table with employees of a company, and information about their immediate supervisors.

- ▶ Write a query which finds the name of the immediate supervisor of all employees
- ▶ Write a query which finds the names of the supervisors of the supervisors of all employees
- ▶ Write a query which finds all the supervisors (transitively) of an employee

Limitations of SQL

- ▶ For a long time, SQL lacked recursion (and loops), and many DBMS's still do
- ▶ That means there are many simple things we can't do easily in SQL, such as traversing a varying number of steps in some kind of list-like structure (e.g., a line of ascendants in a tree)
- ▶ That shortcoming can be dealt with in several ways:
 - ▶ by moving the iterative code to the clients (so we write our recursive calls and loops in another language, like Java or Python)
 - ▶ by using *stored procedures*
 - ▶ by using a graph database
- ▶ Many DBMS's now have some kind of recursion (SQLite3 is one of them), but using it is somewhat difficult and error prone

Recursion in PostgreSQL/SQLite

- ▶ In PostgreSQL and SQLite3, as an example, we can 'loop' to create a table with the numbers 1..10:

```
WITH RECURSIVE cnt(x) AS (  
  VALUES(1)  
  UNION  
  SELECT x+1  
  FROM cnt  
  WHERE x<10  
)  
SELECT x FROM cnt;
```

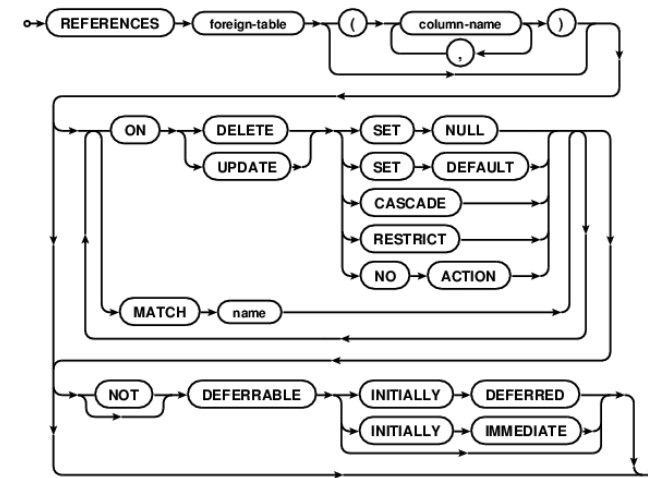
- ▶ A detailed understanding of recursive queries is beyond the scope of the course, but I want you to be aware of them

Foreign key checking

- ▶ When we use foreign key constraints, how do we handle changes in the 'foreign' table?
- ▶ In CREATE TABLE we have several ways to define what should happen when a foreign key changes: on both UPDATE and DELETE we can do:
 - ▶ NO ACTION: doesn't do anything (obviously...)
 - ▶ RESTRICT: makes sure that the parent key can't change when there are foreign keys mapped to them
 - ▶ SET NULL: when a parent key is updated or deleted, the attribute is set to NULL
 - ▶ SET DEFAULT: when a parent key is updated or deleted, the attribute is set to its default value
 - ▶ CASCADE: propagates the change from the parent table to the referencing tables
- ▶ Some of the actions may put the database in a corrupt state
- ▶ In SQLite3 we must enable foreign key checking with:


```
PRAGMA foreign_keys = ON;
```

Foreign key checking



Computer times adjusted to human scale (2017)

One CPU cycle	0.4 ns	1 s
Level 1 cache access	0.9 ns	2 s
Level 2 cache access	2.8 ns	7 s
Level 3 cache access	28 ns	1 min
Main memory access (DDR DIMM)	100 ns	4 min
Intel Optane DC SSD I/O	<10 µs	7 hrs
NVMe SSD I/O	25 µs	17 hrs
SSD I/O	50-150 µs	1.5-4 days
Rotational disk I/O	1-10 ms	1-9 months
Internet call: San Francisco to New York City	65 ms	5 years
Internet call: San Francisco to Hong Kong	141 ms	11 years

Practical consequences of the timescale

- ▶ We try to keep the number of network calls to a minimum – WITH-statements, stored procedures and triggers can help us
- ▶ We try to minimize the number of disk accesses – DBMS's often do this by using B-trees
- ▶ We try to write code which has a memory footprint which is as 'local' as possible (this is not really database related)

Stored procedures

- ▶ We've seen that SQL lacks some simple features, like user defined functions, and loops
- ▶ Often this can be alleviated by writing code in other languages, and have them call SQL queries/statements 'remotely' (like we've done in Java and Python)
- ▶ Some DBMS's have *stored procedures*, which are user defined functions, and in which we can use regular programming features such as parameters and loops
- ▶ One of the main advantages with stored procedures is that it reduces the need for calls over a network connection
- ▶ The way SQLite3 operates makes stored procedures almost pointless – there are no network calls, since the database code is linked into our program

Using a stored procedure in MySQL

```
CREATE FUNCTION getTopLevelSupervisorName(IN emp_name VARCHAR(10))
    RETURN VARCHAR(10)
BEGIN
    DECLARE b_id INT;
    DECLARE b_name VARCHAR(10);

    SET b_name = emp_name;
    SELECT supervisor_id INTO b_id
    FROM employees
    WHERE name = emp_name;
    WHILE b_id <> 0 DO
        SELECT name, supervisor_id
        INTO b_name, b_id
        FROM employees
        WHERE nbr = b_id;
    END WHILE;
    RETURN b_name;
END;
```

Big Data

- ▶ The cost of storing 1 GB of data has gone down from ca SEK 1 000 000 in the 1980ies, to less than SEK 1 today
- ▶ Today, there are many databases on the petabyte scale (that is 1 000 000 000 000 bytes)
- ▶ Data mining is a very active field today, and we mine for:
 - ▶ Scientific discoveries (CERN invented the web for this purpose)
 - ▶ User behavior (customize advertising)
 - ▶ Economic data (algorithmic trading)
 - ▶ ...
- ▶ To find interesting patterns in our data, we often want to store many kinds of data:
 - ▶ text
 - ▶ urls
 - ▶ images
 - ▶ sound
 - ▶ video

SQL vs. NoSQL

- ▶ Relational databases have been a phenomenal success, for several reasons:
 - ▶ ACID
 - ▶ Well tested technology
 - ▶ An enormous amount of money invested in them
- ▶ It is not a panacea for all data, though
 - ▶ They require that we convert all data into tables
 - ▶ They may be more complex than necessary
 - ▶ They're sometimes not fast enough
 - ▶ They don't scale very well (more about that soon)

NoSQL

- ▶ In 2000-2016, many alternatives to relational databases sprung up – for some reason they got the moniker NoSQL (it should really have been NonRelational)
- ▶ Initially it meant literally “No SQL”, but it has evolved into meaning “Not Only SQL”, since many of them embed some SQL like query language themselves
- ▶ NoSQL databases allow users to save and access all sorts of data in simple ways
- ▶ NoSQL databases often sacrifice some of the ACID properties, sometimes they replace “Consistency” with “Eventual consistency”

NoSQL

- ▶ There are different kinds of NoSQL databases:
 - ▶ *Key-value store*: this is the simplest kind of database, it's basically a $\text{Map}[K, V]$, where the database has no way of querying on the contents – Redis is a popular key-value store
 - ▶ *Document store*: this is an enhanced key-value store, where we can search and manipulate data based on the contents, not only the keys – MongoDB and CouchDB are popular document stores
 - ▶ *Column store*: here we have tables with rows containing columns of data, but unlike relational databases, the names and formats of the columns can vary between rows – BigTable and Apache Cassandra are popular column stores
 - ▶ *Graph databases*: Neo4j (from Malmö!) is the most popular example

NoSQL

Exercise

Use MongoDB to insert some student with their applications, and to make some simple queries.

Scaling

- ▶ Handling bigger amounts of data requires that we add computing power and storage space, we can do it in at least two ways:
 - ▶ *Vertical scaling*: updating our processor and expanding our memory – this is also called *scaling up*
 - ▶ *Horizontal scaling*: adding more processors/computers and more hard drives – this is also called *scaling out*
- ▶ Horizontal scaling is normally cheaper, and has greater potential

Capitalizing on horizontal scaling

- ▶ If we scale horizontally, we need a way to use our resources in parallel
- ▶ Google used *map-reduce* for a long time – a simplistic description:
 1. We distribute our data to many servers
 2. On each server, we 'map' a function to our data
 3. We then 'reduce' the results from each server into one final result
- ▶ map and reduce are staples of functional programming
- ▶ Hadoop is an open source framework based on map-reduce



Horizontal scaling of databases

- ▶ Databases are sometimes split row-wise into non-overlapping partitions, this is called *horizontal partitioning*
- ▶ If the partitions are put on separate servers, we call it *sharding*
- ▶ A global company could potentially use sharding to put European customers on European servers, and American customers on American servers – this could make some queries faster
- ▶ Horizontal partitioning (and sharding in particular) enables us to use map-reduce-like algorithms
- ▶ The cost of sharding is increased complexity, and a reliance on the connections between the servers
- ▶ MongoDB and Spanner are examples of DBMS's using sharding

